

Introduction to Deep Learning

Deep Learning and Convolution

Alexandre Allauzen



Jan. 2025


Outline

- 1 Deep Learning: introduction
- 2 Vanishing gradient
- 3 Regularization and normalization
- 4 Convolution for image processing


Outline

- 1 Deep Learning: introduction
- 2 Vanishing gradient
- 3 Regularization and normalization
- 4 Convolution for image processing


Two layers fully connected: a linear separation



Two layers fully connected: a linear separation



Two layers fully connected: a linear separation



Activation f :


- f is usually a non-linear function
 - f is a component wise function
 - tanh, sigmoid, relu, ...
- e.g the softmax function:

Dimensions:

- $\mathbf{x} : D \times 1$
- $\mathbf{W} : M \times D$
- $\mathbf{y} : (M \times \cancel{D}) \times (\cancel{D} \times 1) = M \times 1$

$$y_k = P(c = k | \mathbf{x}) = \frac{e^{\mathbf{w}_k^t \mathbf{x}}}{\sum_{k'} e^{\mathbf{w}_{k'}^t \mathbf{x}}} = \frac{e^{\mathbf{W}_{k,:} \mathbf{x}}}{\sum_{k'} e^{\mathbf{W}_{k',:} \mathbf{x}}}$$

From linear to non-linear case



$$\theta = (\mathbf{W}^{(1)}, \mathbf{W}^{(2)})$$

Trained by
back-propagation of
the gradient

Universal approximation theorem



a feed-forward network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of \mathbb{R}^n , under mild assumptions on the activation function. (...)

(Cybenko1989)

However, it does not touch upon the algorithmic learnability of those parameters.




Overfitting

The danger of the over-parametrization





Source: Wikipedia



A first example





A first example, with 3 hidden units




The xor-like example





The xor-like example, with 3 hidden units




A more difficult example



A more difficult example, with 3 hidden units



A more difficult example, with 8 hidden units

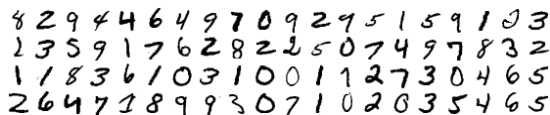


Outline

- 1 Deep Learning: introduction
- 2 Vanishing gradient
- 3 Regularization and normalization
- 4 Convolution for image processing


Experimental observations (MNIST task) - 1

The MNIST database



8 2 9 4 4 6 4 9 7 0 9 2 9 5 1 5 9 1 0 3
 1 3 5 9 1 7 6 2 8 2 2 5 0 7 4 9 7 8 3 2
 1 1 8 3 6 1 0 3 1 0 0 1 1 2 7 3 0 4 6 5
 2 6 4 7 1 8 9 9 3 0 7 1 0 2 0 3 5 4 6 5

Comparison of different depth for feed-forward architecture




- Hidden layers have a sigmoid activation function.
- The output layer is a softmax.

Experimental observations (MNIST task) - 2

Varying the depth


- Without hidden layer: $\approx 88\%$ accuracy
- 1 hidden layer (30): $\approx 96.5\%$ accuracy
- 2 hidden layers (30): $\approx 96.9\%$ accuracy
- 3 hidden layers (30): $\approx 96.5\%$ accuracy
- 4 hidden layers (30): $\approx 96.5\%$ accuracy



(From <http://neuralnetworksanddeeplearning.com/chap5.html>)

Intuitive explanation

Let consider the simplest deep neural network, with just a single neuron in each layer.



w_i, b_i are resp. the weight and bias of neuron i and C some cost function.


Compute the gradient of C *w.r.t* the bias b_1

$$\frac{\partial C}{\partial b_1} = \frac{\partial C}{\partial y_4} \times \frac{\partial y_4}{\partial a_4} \times \frac{\partial a_4}{\partial y_3} \times \frac{\partial y_3}{\partial a_3} \times \frac{\partial a_3}{\partial y_2} \times \frac{\partial y_2}{\partial a_2} \times \frac{\partial a_2}{\partial y_1} \times \frac{\partial y_1}{\partial a_1} \times \frac{\partial a_1}{\partial b_1} \quad (1)$$

$$= \frac{\partial C}{\partial y_4} \times \sigma'(a_4) \times w_4 \times \sigma'(a_3) \times w_3 \times \sigma'(a_2) \times w_2 \times \sigma'(a_1) \quad (2)$$

Intuitive explanation - 2

The derivative of the activation function: σ'



$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$


But weights are initialize around 0.

The different layers in our deep network are learning at vastly different speeds:

- when later layers in the network are learning well,
- early layers often get stuck during training, learning almost nothing at all.

A first Solution

Change the activation function (Rectified Linear Unit or ReLU)



- Avoid the vanishing gradient
- Some units can "die"

See (Glorot et al.2011) for more details

Variants

- Leaky ReLU (Maas et al.2013)
- Soft-plus $\log(1 + e^x)$

And many more, see <https://pytorch.org/docs/stable/nn.html>

More details

See (Hochreiter et al.2001; Glorot and Bengio2010; LeCun et al.2012)

A question

Why adding a layer can lower the performance ?


- Overfitting ? and what about the identity
- Vanishing gradient ? and with the *Relu* ?

Residual block

From (He et al.2016)

- Add a skip connection
- The model learn the "residual"

$$\mathbf{y} = \mathbf{x} + \mathcal{F}(\mathbf{x})$$



A simple version of highway networks (Srivastava et al.2015)

Residual block

Forward

$$\mathbf{y} = \mathbf{x} + \mathcal{F}(\mathbf{x})$$


Backward

Assume a residual block for the layer l in the network. Training requires:

- $\frac{\partial l}{\partial \mathbf{W}^{(l)}}$ for the update of the layer
- $\frac{\partial l}{\partial \mathbf{x}^{(l)}}$ for the backpropagation

$$\begin{aligned} \frac{\partial l}{\partial \mathbf{x}^{(l)}} &= \frac{\partial l}{\partial \mathbf{y}^{(l)}} \times \frac{\partial \mathbf{y}^{(l)}}{\partial \mathbf{x}^{(l)}} \\ &= \frac{\partial l}{\partial \mathbf{y}^{(l)}} \times \left(1 + \frac{\partial \mathcal{F}(\mathbf{x}^{(l)})}{\partial \mathbf{x}^{(l)}}\right) \end{aligned}$$

ResNet in action



<http://kaiminghe.com/>

Outline

- 1 Deep Learning: introduction
- 2 Vanishing gradient
- 3 Regularization and normalization**
- 4 Convolution for image processing

Regularization l^2 or gaussian prior or weight decay

The basic way:


$$\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) = \sum_{i=1}^N l(\boldsymbol{\theta}, \mathbf{x}_{(i)}, c_{(i)}) + \frac{\lambda}{2} \|\boldsymbol{\theta}\|^2$$

- The second term is the **regularization term**.
- Each parameter has a gaussian prior : $\mathcal{N}(0, 1/\lambda)$.
- λ is a hyperparameter.
- The update has the form:


$$\boldsymbol{\theta} = (1 - \eta_t \lambda) \boldsymbol{\theta} - \eta_t \nabla_{\boldsymbol{\theta}}$$

Dropout

A new regularization scheme (Srivastava and Salakhutdinov 2014)



(a) Standard Neural Net



(b) After applying dropout.

- For each training example: randomly turn-off the neurons of hidden units (with $p = 0.5$)
- At test time, use each neuron scaled down by p

- Dropout serves to separate effects from strongly correlated features and
- prevents co-adaptation between units
- It can be seen as averaging different models that share parameters.
- It acts as a powerful regularization scheme.

Dropout - implementation

The layer should keep:

- $\mathbf{W}^{(l)}$: the parameters
- $f^{(l)}$: its activation function
- $\mathbf{x}^{(l)}$: its input
- $\mathbf{a}^{(l)}$: its pre-activation associated to the input
- $\delta^{(l)}$: for the update and the back-propagation to the layer $l - 1$
- $\mathbf{m}^{(l)}$: the dropout mask, to be applied on $\mathbf{x}^{(l)}$

Forward pass

For $l = 1$ to $(L - 1)$

- Compute $\mathbf{y}^{(l)} = f^{(l)}(\mathbf{W}^{(l)}\mathbf{x}^{(l)})$
- $\mathbf{x}^{(l+1)} = \mathbf{y}^{(l)} = \mathbf{y}^{(l)} \circ \mathbf{m}^{(l)}$

$$\mathbf{y}^{(L)} = f^{(L)}(\mathbf{W}^{(L)}\mathbf{x}^{(L)})$$

Internal Covariate Shift

After init.: training

For a layer l ,

- The update of $\mathbf{W}^{(l)}$ depends on the incoming gradient (from the loss)
- And a mini-batch of $\mathbf{X}^{(l)} = (\mathbf{x}^{(l)})$, but the statistics $\mathbf{X}^{(l)}$ depends on $\mathbf{W}^{(l-1)}$
- During the back-propagation, you then update $\mathbf{W}^{(l-1)}$.
- During the next forward-backward pass, the statistics of $\mathbf{X}^{(l)}$ have changed !
- It could have the same effect as a bad init. !

Sketch of solution

- Stabilize the distribution of each $\mathbf{X}^{(l)}$.
- By learning a layer that make each dimension zero-mean unit-variance.

Batch normalization

Assume a mini batch of dim. (B, D)

Compute statistics and normalization

$$\mu_j = \frac{1}{B} \sum_{i=1}^B x_{i,j} \quad \text{and} \quad \sigma_j^2 = \frac{1}{B} \sum_{i=1}^B (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Record also running average and variance for the test time.

Shift and scale

$$x_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

With γ and β are the learnable vectors of the batch-norm layer.

See (Ioffe and Szegedy2015)

How to apply batch norm

Where to insert the batch-norm

$$\mathbf{W}_1 \rightarrow \textcolor{red}{BN}_1 \rightarrow \textit{ReLU} \rightarrow \mathbf{W}_2 \rightarrow \textcolor{red}{BN}_2 \rightarrow \textit{ReLU} \rightarrow \dots$$

but this one is better:

$$\mathbf{W}_1 \rightarrow \textit{ReLU} \rightarrow \textcolor{red}{BN}_1 \rightarrow \mathbf{W}_2 \rightarrow \textit{ReLU} \rightarrow \textcolor{red}{BN}_2 \rightarrow \dots$$

Drawbacks

- The batch size must be large enough to capture overall statistics, which is sometimes impossible if you are working with large models.
- The concept of a batch is not always present, or it may change from time to time.
- Not the same behaviour in train and test

Layer Norm

Independant of the batch size with no difference between train and test.

Compute statistics and normalization per dimension

$$\mu_i = \frac{1}{D} \sum_{j=1}^D x_{i,j} \quad \text{and} \quad \sigma_i^2 = \frac{1}{D} \sum_{j=1}^D (x_{i,j} - \mu_i)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}$$

$$x_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

With γ and β are the learnable vectors of the batch-norm layer.

Works well also with more advanced architectures (recurrent and transformers)

See (Ba et al.2016)

Outline

- 1 Deep Learning: introduction
- 2 Vanishing gradient
- 3 Regularization and normalization
- 4 Convolution for image processing

Convolution for image processing

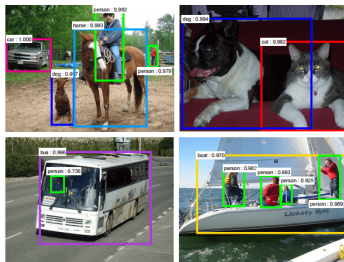


Image classification

An image = (2D array of values) \times (number of channels)

2D array

The spatial structure:

- A 2D real space
- With distance

Channels



- For image : R,G,B
- In fluid mechanics: pression, velocity, ...
- In general: different measures on the same spatial domain

Sources



Many figures and examples are inspired by, or extracted from the course of the Stanford course of Fei-Fei Li.

<http://cs231n.stanford.edu/>

Image classification: the flatten way



Another example with a color image (3 channels) of 32×32 pixels



Feed-forward NNet for image classification

Question ?

- Assume a NNet for images of size 256×256 in color,
 - with one hidden layer of 1000 and an output layer for 100 classes
- How many parameters do we have ?
- And for a larger image like a 16 : 9 resolution (1920×1080) ?

Intuition for convolution

- Look at a small region (window)
- Apply a (small) filter
- Translation invariance by sliding the window along the image
- Parameters sharing

Convolution 2D - the first step

Extract a frame, or a window, and apply a “filter”

The input image ($L = 6, C = 6$):
and the window

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

Apply the filter:
kernel size of $ks = (3, 3)$

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$

The output value (output channel)

$$\text{First value: } h_{1,1} = \sum_{i,j} w_{i,j} \times x_{i,j}$$

Convolution 2D - slide the window

"Scan" the image with the filter

The **next** window

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

Apply the **same** filter
with a **stride** (1,1)

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$

The output value (output channel)

$$\text{Second value: } h_{1,2} = \sum_{i,j} w_{i,j} \times x_{i,j}$$

$$\text{In general: } h_{l,c} = \sum_{i,j=(0,0)}^{(I,J)} w_{i,j} \times x_{l+i,c+j}$$

2D Convolution: animated version (1st row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
-----------	-----------	-----------	-----------

2D Convolution: animated version (1st row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
-----------	-----------	-----------	-----------

2D Convolution: animated version (1st row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
-----------	-----------	-----------	-----------

2D Convolution: animated version (1st row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
-----------	-----------	-----------	-----------

2D Convolution: animated version (2nd row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
$h_{2,1}$	$h_{2,2}$	$h_{2,3}$	$h_{2,4}$

2D Convolution: animated version (2nd row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
$h_{2,1}$	$h_{2,2}$	$h_{2,3}$	$h_{2,4}$

2D Convolution: animated version (2nd row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
$h_{2,1}$	$h_{2,2}$	$h_{2,3}$	$h_{2,4}$

2D Convolution: animated version (2nd row)

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$	$h_{1,4}$
$h_{2,1}$	$h_{2,2}$	$h_{2,3}$	$h_{2,4}$

2D Convolution: the final result

With a stride of (1, 1)

The input image:

$x_{1,1}$	$x_{2,1}$	$x_{3,1}$	$x_{4,1}$	$x_{5,1}$	$x_{6,1}$
$x_{1,2}$	$x_{2,2}$	$x_{3,2}$	$x_{4,2}$	$x_{5,2}$	$x_{6,2}$
$x_{1,3}$	$x_{2,3}$	$x_{3,3}$	$x_{4,3}$	$x_{5,3}$	$x_{6,3}$
$x_{1,4}$	$x_{2,4}$	$x_{3,4}$	$x_{4,4}$	$x_{5,4}$	$x_{6,4}$
$x_{1,5}$	$x_{2,5}$	$x_{3,5}$	$x_{4,5}$	$x_{5,5}$	$x_{6,5}$
$x_{1,6}$	$x_{2,6}$	$x_{3,6}$	$x_{4,6}$	$x_{5,6}$	$x_{6,6}$

The filter:

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$



$h_{1,1}$	$h_{2,1}$	$h_{3,1}$	$h_{4,1}$
$h_{1,2}$	$h_{2,2}$	$h_{3,2}$	$h_{4,2}$
$h_{1,3}$	$h_{2,3}$	$h_{3,3}$	$h_{4,3}$
$h_{1,4}$	$h_{2,4}$	$h_{3,4}$	$h_{4,4}$

Exercise

Stride

- With an input image of length (L, C) , a kernel size of ks and a stride $= (1, 1)$, what is the output dimension ?
- And with a stride $= (2, 2)$, what is the output dimension ?
- With a stride of s ?

Side effect

- The first and last rows and columns are not processed as the others. How to correct this aspects ?
- How to ensure the same dimension in output (assuming a stride of 1) ?
- How to ensure that every inputs are “seen” equally ?

Padding

Constant padding of (1,2) with c


			c	c	c	c	c	c	
1	4	3	c	c	1	4	3	c	c
0	6	5	c	c	0	6	5	c	c
2	7	9	c	c	2	7	9	c	c
			c	c	c	c	c	c	

Reflection padding of (1,2)

			5	6	0	6	5	6	?
1	4	3	3	4	1	4	3	4	1
0	6	5	5	6	0	6	5	6	0
2	7	9	9	7	2	7	9	7	2
			5	6	0	6	5	6	?

Convolution in 2D

The basics



$3 \times 32 \times 32$ image




$3 \times 5 \times 5$ filter

Convolution of the filter with the image:

- Sliding the filter along the two axis
 - Computing the “dot product” at each step
- preserve the spatial structure along the channels
- each step extract a “local” and spatial feature


Convolution in 2D

With a single output channel




Convolution in 2D

Add a second output channel



Shape of things


- An image is (C_{in}, L, C)



- A convolution filter is (C_{in}, W, V)




- A convolution filter with (C_{out}, C_{in}, W, V)




Motivation for convolution

Extract “low level” features




Motivation for convolution

Extract “low level” features



Motivation for convolution


Extract “low level” features



Max-pooling or Downsampling in 2D


The goal

- Convolution extracts local features (followed by non-linearity)
- Max-pooling acts as a selection, compression, or contraction operator
- The back-propagation promote feature saliency for each channel




Architecture of deep convolution NNet for image processing

VGGNet (Simonyan and Zisserman 2015)



Conv. layers with kernel size of 3×3 , stride and padding, followed by pooling layers (max on 2×2 with stride 2).

ResNet for Imagenet



(a) original

(b) proposed



Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton.
2016.

Layer normalization.



G. Cybenko.

1989.

Approximation by superpositions of a sigmoidal function.

Mathematics of Control, Signals, and Systems (MCSS), 2(4):303–314, December.



Xavier Glorot and Yoshua Bengio.

2010.

Understanding the difficulty of training deep feedforward neural networks.

In *JMLR W&CP: Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics (AISTATS 2010)*, volume 9, pages 249–256, May.



Xavier Glorot, Antoine Bordes, and Yoshua Bengio.

2011.

Deep sparse rectifier neural networks.

In Geoffrey J. Gordon and David B. Dunson, editors, *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS-11)*, volume 15, pages 315–323. Journal of Machine Learning Research - Workshop and Conference Proceedings.



Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun.

2016.

Deep residual learning for image recognition.

In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778.



S. Hochreiter, Y. Bengio, P. Frasconi, and J. Schmidhuber.

2001.

Gradient flow in recurrent nets: the difficulty of learning long-term dependencies.

In Kremer and Kolen, editors, *A Field Guide to Dynamical Recurrent Neural Networks*. IEEE Press.



Sergey Ioffe and Christian Szegedy.

2015.

Batch normalization: Accelerating deep network training by reducing internal covariate shift.

In Francis Bach and David Blei, editors, *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pages 448–456, Lille, France, 07–09 Jul. PMLR.



Yann LeCun, Léon Bottou, Genevieve Orr, and Klaus-Robert Müller.

2012.

Efficient backprop.

In Grégoire Montavon, GenevièveB. Orr, and Klaus-Robert Müller, editors, *Neural Networks: Tricks of the Trade*, volume 7700 of *Lecture Notes in Computer Science*, pages 9–48. Springer Berlin Heidelberg.



Andrew L. Maas, Awni Y. Hannun, and Andrew Y. Ng.

2013.

Rectifier nonlinearities improve neural network acoustic models.

In *in ICML Workshop on Deep Learning for Audio, Speech and Language Processing*.



Karen Simonyan and Andrew Zisserman.

2015.

Very deep convolutional networks for large-scale image recognition.

In *International Conference on Learning Representations*.



Nitish Srivastava and Ruslan Salakhutdinov.

2014.

Multimodal learning with deep boltzmann machines.

Journal of Machine Learning Research, 15:2949–2980.



Rupesh Kumar Srivastava, Klaus Greff, and Jürgen Schmidhuber.

2015.

Training very deep networks.

CoRR, abs/1507.06228.