

TITLE: Text, word embeddings, transformers

Roadmap

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Text classification/rating

my wonderful friend took
me to see this movie
for our anniversary.
it was terrible.

☹️ : 0 ... 1 : 😊

- How to represent the input text ?
- How to make classification ?

Bag of words (BOW)

this movie is just great , with a great music , while a bit long

Bag of words (BOW)

this movie is just great , with a great music , while a bit long

vocabulary	binary bag	count bag	tf.idf bag	...
awesome	0	0	0	...
great	1	2	1.9	...
long	1	1	2.5	...
the	0	0	0	...
this	1	1	0.1	...

A basic vectorial representation of text

$$\mathbf{x} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 0 \\ 1 \end{pmatrix} \in \mathbb{R}^D$$

$$\left. \begin{array}{l} \textit{awesome} \\ \textit{great} \\ \textit{long} \\ \textit{the} \\ \textit{this} \end{array} \right\} D$$

A simple problem

Assumptions

- Let define a finite set of known words: the vocabulary \mathcal{V}
- A text is a vector \mathbf{x} of dimension $D = |\mathcal{V}|$
- Each component encodes the presence of a word

Then machine learning

- Naive Bayes
- SVM, Random Forrest, ...
- Logistic Regression

Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Back to logistic regression

$$\mathbf{x} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 0 \\ 1 \end{pmatrix} \in \mathbb{R}^D \quad \left. \begin{array}{l} \text{awesome} \\ \text{great} \\ \text{long} \\ \text{the} \\ \text{this} \end{array} \right\} D$$

For one input text:

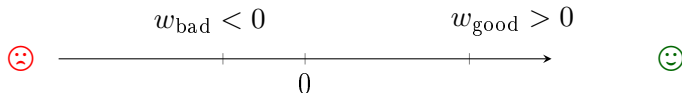
$$w_0 + \mathbf{w}^t \mathbf{x} = w_0 + 2 \times w_2 + w_3 + w_5$$

The class is positive ($y = 1$) if

$$\begin{aligned} w_0 + 2 \times w_2 + w_3 + w_5 &> 0 \\ 2 \times w_{\text{great}} + w_{\text{long}} + w_{\text{this}} + &> -w_0 \end{aligned}$$

A limited representation of words

With the logistic regression model on a bag of words:



Consider the two following examples:

the end is really bad		☹️ $\Rightarrow w_{\text{bad}}$	\searrow
the bad guy is <i>awesome</i>		😊 $\Rightarrow w_{\text{bad}}$	$\searrow, w_{\text{awesome}} \nearrow$

Multiple dimensions could help to:

- represent different usage
- consider the context
- leverage more from sparse, sometime ambiguous observations.

A simple model for document classification - part 1

Idea

- The word representation could be shared among classes
- While their interpretation depends on the class

Input representation and composition

$$\mathbf{R} \times \mathbf{x} = \begin{pmatrix} \vdots & \vdots & \vdots & \vdots & \vdots \\ \mathbf{v}_1 & \mathbf{v}_2 & \mathbf{v}_3 & \mathbf{v}_4 & \mathbf{v}_5 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix} \times \begin{pmatrix} 0 \\ \mathbf{2} \\ \mathbf{1} \\ 0 \\ \mathbf{1} \end{pmatrix} = 2 \times \mathbf{v}_2 + \mathbf{v}_3 + \mathbf{v}_5 = \mathbf{d}$$

A simple model for document classification - part 2

Classification

$$\begin{aligned}P(y|\mathbf{x}) &= \text{softmax}(\mathbf{W}^o \mathbf{d}) = \text{softmax}(\mathbf{W}^o \times \mathbf{R}\mathbf{x}), \text{ or} \\ &= \text{softmax}(\mathbf{W}^o \times f(\mathbf{R}\mathbf{x})),\end{aligned}$$

with f a non-linear activation function.

Parameters

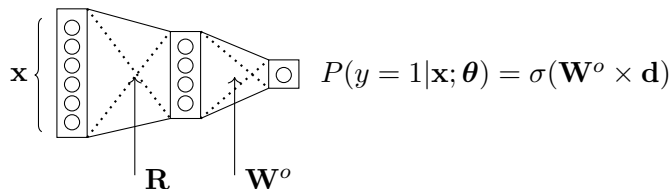
$$\theta = (\mathbf{R}, \mathbf{W}^o) \rightarrow \text{to learn !!}$$

Reminder

If $\mathbf{y} = \text{softmax}(\mathbf{a})$, \mathbf{y} is a vector and \mathbf{a} is called the logit vector

$$y_i = \frac{e^{a_i}}{\sum_j e^{a_j}}$$

A first neural network



- $\mathbf{x} : (|\mathcal{V}|, 1)$
- $\mathbf{R} : (K, |\mathcal{V}|)$
- $\mathbf{d} : (K, 1)$
- $\mathbf{W}^o : (1, K)$
- $y : (1, 1)$

$$\mathbf{d} = \mathbf{R} \times \mathbf{x}$$

$$y = \sigma(\mathbf{W}^o \times \mathbf{d})$$

Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Limitations of BOW classifier

$$\mathbf{h} = \sum_{i=1}^L \underbrace{\mathbf{x}_i}_{\text{emb. of word } i}$$

Limitations

- Words are equally important
- Word order independent
- Miss contextual information (local/global)

Local contexts

the	end	is	very	bad	but	what	a	great	music
-----	-----	----	------	-----	-----	------	---	-------	-------

Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad ++$						

Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad++$						
			$\underbrace{\hspace{2.5cm}}$ $but \text{ will change } bad$						

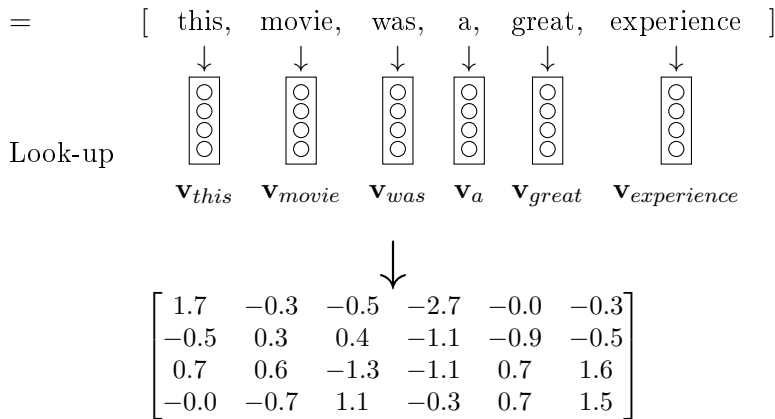
Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$						
			<i>very</i> \rightarrow <i>bad</i> ++						
			$\underbrace{\hspace{2.5cm}}$						
						<i>but</i> will change <i>bad</i>			
		$\underbrace{\hspace{2.5cm}}$						$\underbrace{\hspace{2.5cm}}$	
		<i>bad</i> is for <i>end</i> not <i>music</i>							
								<i>great</i> is for <i>music</i> not fo <i>end</i>	

Motivations

- Local contextualisation
- Global view of the sentence

Another view of a sentence



Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Draw attention for classification

Remind CBOW classifier

The classifier output:

$$\text{softmax}(\mathbf{W}^o \mathbf{h}) \text{ (multiclass) or } \sigma(\mathbf{w}^o \mathbf{h}) \text{ (binary)}$$

- What does represent a row of \mathbf{W}^o ?
- The product $\mathbf{W}^o \mathbf{h}$?
- The softmax ?

Draw attention

Is a word vector related to the classification task ?

$$\mathbf{h} = \sum_{i=1}^L \underbrace{\mathbf{x}_i}_{\text{emb. of word } i} \longrightarrow \mathbf{h} = \sum_{i=1}^L \underbrace{\lambda_i}_{\text{???}} \mathbf{x}_i$$

Draw attention for classification (binary task)

$$\mathbf{X}\mathbf{q} = L \left\{ \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline \end{array} \right\} \times \begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array} = \begin{array}{|c|} \hline \\ \hline \\ \hline \\ \hline \\ \hline \\ \hline \end{array} \in \mathbb{R}^L$$
$$(\mathbf{X}\mathbf{q})_i = \mathbf{x}_i^t \mathbf{q} \quad (\text{dot product})$$
$$\mathbf{a} = \text{softmax}(\mathbf{X}\mathbf{q})$$

- $\mathbf{a} = (a_i)$, $\sum_{i=1}^L a_i = 1$ and $0 \leq a_i \leq 1$
- \mathbf{a} : attention vector for the "query" \mathbf{q} and the "keys" \mathbf{X} .
- \mathbf{q} is a vector to be learnt [11, 7]

Attention to weight inputs (binary task)

- $\mathbf{a} = \text{softmax}(\mathbf{X}\mathbf{q})$ is the attention vector

$$\mathbf{h} = \sum_{i=1}^L a_i \mathbf{x}_i = \mathbf{a}^t \mathbf{X}$$


- A new vector, focused on the classification task (\mathbf{q})
- To summarize:

$$\mathbf{h} = \text{softmax}(\mathbf{X}\mathbf{q})^t \mathbf{X} \rightarrow \text{classification}$$

Issues:

- Scale the dot product
- \mathbf{X} is involved everywhere !

Basic attention mechanism for classification (binary task)



$$\mathbf{h} = \text{softmax} \left(\frac{\mathbf{K}\mathbf{q}}{\sqrt{d}} \right)^t \mathbf{V}$$

- \mathbf{X} can be static emb.
- or **contextualized embedding**
- \mathbf{q} is learnt as a target for selection
- $\mathbf{a} = \mathbf{K}\mathbf{q}$: selection in \mathbf{V}

Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Contextualized word embeddings

Consider the word **driver**:

the audio **driver** is really outdated
the **driver** exceeded the speed limit

The context

The	■	The	■	$\lambda_{2,1}$
audio	■■■	driver	■■■■■	$\lambda_{2,2}$
driver	■■■■■	exceeded	■	$\lambda_{2,3}$
is	■	the	■	$\lambda_{2,4}$
really	■	speed	■■■	$\lambda_{2,5}$
outdated	■■■	limit	■	$\lambda_{2,6}$

Self attention: a first idea

Look at the "correlation" between words (embeddings)

- $\mathbf{X}\mathbf{X}^t$ is a $L \times L$ matrix, stores $(\mathbf{x}_i^t \mathbf{x}_j)$
- The i^{th} row stores the "correlation between" \mathbf{x}_i and all the other words in the sentence
- For $i = 2$, we have the correlations with **driver**
- We can use this correlation as a weight

$$\mathbf{z}_2 = \mathbf{z}_{driver} = \sum_{j=1}^L \underbrace{\lambda_{2,j}}_{\mathbf{x}_2^t \mathbf{x}_j} \mathbf{x}_j$$

More (linear) transformations

Two different Transformations on \mathbf{X}

$$\mathbf{X} \longrightarrow \mathbf{X}\mathbf{W}_Q = \mathbf{Q}$$

$$\mathbf{X} \longrightarrow \mathbf{X}\mathbf{W}_K = \mathbf{K},$$


- with \mathbf{W}_Q and $\mathbf{W}_K \in \mathbb{R}^{d \times d}$
- \mathbf{Q} and \mathbf{K} have the same dimensions as \mathbf{X}

$$\mathbf{A} = \mathbf{Q}\mathbf{K}^t = \underbrace{(\mathbf{Q}_{i,*}\mathbf{K}_{j,*}^t)_{i,j}}_{L \times L} = (\mathbf{q}_i^{\mathbf{k}^j}) = (\lambda_{i,j}),$$

with $\lambda_{i,j}$ the attention on "word" j to generate \mathbf{z}_i

Normalization of attention

Take the row-wise softmax:

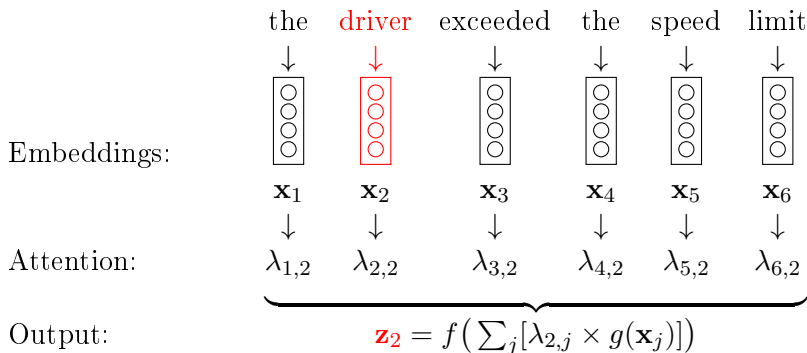


$$\sum_j \underbrace{\lambda_{i,j}}_{\text{or } a_{i,j}} = 1 \text{ and } \lambda_{i,j} \geq 0$$

Each row of \mathbf{A} gives a convex combination

Self attention (overview)


Consider the word **driver**:



- $(\lambda_{i,j})$ are the attention coefficients, $\sum_j \lambda_{i,j} = 1$, and
- Reflects the influence of \mathbf{x}_j on \mathbf{x}_i (transformed version)


Transformer : Queries, Keys, Values

the driver exceeded the speed limit

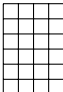


Tranformer : Attention matrix


The distance matrix between Q and K



Scaled Dot-Product Attention

$$\mathbf{Z} = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^t}{\sqrt{d}}\right)\mathbf{V} =$$


Multi-head attention (with 2 heads)




Putting all together (with more tricks)

Transformer block

From [10]


- Inputs is \mathbf{X}
- Positional embeddings
- Multihead attention
- Residual connections [6]
- Layer Normalization [2]
- Final filtering



Layer norm


Assume \mathbf{Z} a minibatch of sequences (B, L, D) : $\mathbf{Z} = L \left\{ \begin{array}{c} \text{grid} \\ \vdots \\ \text{grid} \end{array} \right\}$
 d

Batch or Layer norm



[9]

Positional embeddings




- Originally "absolute"
- Can be learnt [5, 1]
- Or relative [8]

(figure generated by the following code


```
https://github.com/jalammar/jalammar.github.io/blob/master/notebooks/  
transformer/transformer\_positional\_encoding\_graph.ipynb)
```

A Transformer layer

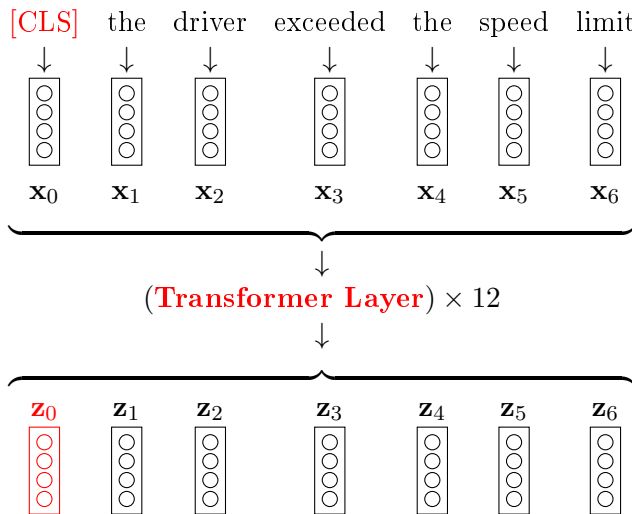


Transformer layers can be stacked !

Pre-training as a (Masked) language model



BERT Encoder for text classification



Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

Summary

Attention, attention

- This mechanism allows the model to efficiently handle different kind of structure.
- Originally for machine translation, and with BI-GRU [4, 3].

Transformers

- Architecture proposed in [10]
- Nowadays state of the art component

Transformers are everywhere

State of the art encoder

- For text ! (BERT)
- And also for speech, DNA, vision, ...

Also a powerful generator

- For text (GPT, ...)
- Speech, ... sequences

Outline

Introduction

Word embeddings

Text classification, beyond BOW

Attention for classification

Transformer architecture

Conclusion

References

- [1] Rami Al-Rfou et al. *Character-Level Language Modeling with Deeper Self-Attention*. 2018. arXiv: 1808.04444 [cs.CL].
- [2] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. *Layer Normalization*. 2016. arXiv: 1607.06450 [stat.ML].
- [3] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. “Neural Machine Translation by Jointly Learning to Align and Translate”. In: *CoRR* abs/1409.0473 (2014). URL: <http://arxiv.org/abs/1409.0473>.
- [4] Kyunghyun Cho et al. “Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation”. In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Doha, Qatar: Association for Computational Linguistics, 2014, pp. 1724–1734. URL: <http://www.aclweb.org/anthology/D14-1179>.
- [5] Jonas Gehring et al. “Convolutional Sequence to Sequence Learning”. In: *CoRR* abs/1705.03122 (2017). arXiv: 1705.03122. URL: <http://arxiv.org/abs/1705.03122>.
- [6] Kaiming He et al. “Deep Residual Learning for Image Recognition”. In: *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 770–778. arXiv: 1512.03385 [cs.CV].
- [7] Zhouhan Lin et al. “A STRUCTURED SELF-ATTENTIVE SENTENCE EMBEDDING”. In: *International Conference on Learning Representations*. 2017. URL: https://openreview.net/forum?id=BJC_jUqxe.

- [8] Peter Shaw, Jakob Uszkoreit, and Ashish Vaswani. *Self-Attention with Relative Position Representations*. 2018. arXiv: 1803.02155 [cs.CL].
- [9] Sheng Shen et al. “PowerNorm: Rethinking Batch Normalization in Transformers”. In: *Proceedings of the 37th International Conference on Machine Learning*. Ed. by Hal Daumé III and Aarti Singh. Vol. 119. Proceedings of Machine Learning Research. PMLR, 2020, pp. 8741–8751. URL: <https://proceedings.mlr.press/v119/shen20e.html>.
- [10] Ashish Vaswani et al. “Attention is All you Need”. In: *Advances in Neural Information Processing Systems 30*. Ed. by I. Guyon et al. Curran Associates, Inc., 2017, pp. 6000–6010. URL: <http://papers.nips.cc/paper/7181-attention-is-all-you-need.pdf>.
- [11] Zichao Yang et al. “Hierarchical Attention Networks for Document Classification”. In: *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)06*. 2016.